

# 225 homework 3

Jing Liao

Consider the stroke diagnostic problem from Assignment 2. This time, use a Gaussian process classification model to predict stroke.

For categorical outcome variable, we assume the Gaussian process prior over a continuous latent function,  $u(x)$ . We define the distribution of the response variable in terms of this latent function. The outcome variable in our problem is binary, I would use the following logistic model

$$P(y_i = 1 | u(x_i)) = \frac{\exp(u(x_i))}{1 + \exp(u(x_i))}$$

We need to choose an appropriate covariance function in setting Gaussian process models.

Firstly, I tried Squared exponential kernel, this kernel is most flexible kernel compared to OU and Brownian motion.

$$C_{ij} = \sigma^2 \exp\left[-\sum_{p=1}^p \rho_p (X_{ip} - X_{jp})^2\right]$$

We have 14 predictors in our model, there is one  $\rho$  for each predictor, where  $\sigma \sim U(0.1, 1)$  and  $\rho_i \sim Ga(5, 5)$

```
library(R2jags)
stroke=read.table("/Users/jing/Desktop/2020 winter/Stat 225/hw2/Stroke.csv",sep=',',header=T)
stroke$Gender<-as.numeric(stroke$Gender)-1
attach(stroke)
N = dim(stroke)[1];
p = dim(stroke)[2]
Y=stroke[,16]
X = as.matrix(stroke[,2:(p-1)])
#x=scale(x[,2:14])
#x = cbind(stroke$Gender, x)
tset1 <- Sys.time()
gpsel <- "model{

  # Likelihood

  for(i in 1:T){
    Y[i] ~ dbern(p[i])
    p[i]<-exp(u[i])/(1+exp(u[i]))
  }
  u~ dmnorm(Mu, Sigma.inv)
  Sigma.inv <- inverse(Sigma)

  # Set up mean and covariance matrix

  for(i in 1:T) {
    Mu[i] <- alpha
    Sigma[i,i] <- pow(sigma, 2)
    for(j in (i+1):T) {
      for (p in 1:14){
```

```

      C1[i,j,p]=-(rho[p] * pow(X[i,p] - X[j,p], 2))
      #C2[i,j,p]= pow(X[i,p] - X[j,p], 2)
    }
    Sigma[i,j] <- exp(sum(C1[i,j,]))
    Sigma[j,i] <- Sigma[i,j]
  }
}
#Priors

alpha ~ dnorm(0, 0.01)
sigma ~ dunif(0, 1)
for (p in 1:14){
rho[p] ~ dgamma(5, 5)
}

}"

dat<- list(Y=Y,T=100,X=X)
jags.param=c('alpha','rho','sigma')
fitse1 <- jags(data=dat, n.chains=4, inits=NULL,parameters=jags.param, n.iter=3000, n.burnin=1000,
               model.file=textConnection(gpse1))
tse2 <- Sys.time()
tse2-tse1
print(jagsfitse1)

```

Since this model is flexible and takes a long time to run. It take about 1.5 hours to run with DIC 137.6.

Then I tried the OU kernel,

$$C_{ij} = \sigma^2 \exp\left[-\sum_{p=1}^p \frac{1}{\rho} |X_{ip} - X_{jp}|\right]$$

where  $\sigma \sim U(0.1, 1)$  and  $\rho_i \sim U(1, 2)$

```

to1 <- Sys.time()
gpou <- "model{

  # Likelihood

  for(i in 1:T){
    Y[i] ~ dbern(p[i])
    p[i]<-exp(u[i])/(1+exp(u[i]))
  }
  u~ dmnorm(Mu, Sigma.inv)
  Sigma.inv <- inverse(Sigma)

  # Set up mean and covariance matrix

  for(i in 1:T) {
    Mu[i] <- alpha
    Sigma[i,i] <- pow(sigma, 2)
    for(j in (i+1):T) {
      for (p in 1:14){
        C[i,j,p]=(1/rho[p])*abs(X[i,p]-X[j,p])
      }
      Sigma[i,j] <- pow(sigma, 2) * exp(-sum(C[i,j,]))
    }
  }
}

```

```

        Sigma[j,i] <- Sigma[i,j]
    }

}
#Priors

alpha ~ dnorm(0, 0.01)
sigma ~ dunif(0.1, 1)
for (p in 1:14){
    rho[p] ~ dunif(1, 2)
}

}"

dat<- list(Y=Y,T=100,X=X)
jags.param=c('alpha','sigma','p')
fitou <- jags(data=dat, n.chains=4, inits=NULL,parameters=jags.param, n.iter=3000, n.burnin=1000,
              model.file=textConnection(gpou))
to2 <- Sys.time()
to2-to1
print(fitou)

```

The OU kernel Gaussian process takes about 1 hours to run with DIC 135 which it is time consuming.

Since Brownian motion is still simpler than the OU process, the computational cost of inverting  $C$  is  $O(n)$

$$C_{ij} = \sum_{p=1}^p \sigma_p^2 \min(X_{ip}, X_{jp})$$

where  $\sigma_p \sim U(1, 2)$

```

gpbm1 <- "model{

    # Likelihood

    for(i in 1:T){
        Y[i] ~ dbern(p[i])
        p[i]<-exp(u[i])/(1+exp(u[i]))
    }
    u~ dmnorm(Mu, Sigma.inv)
    Sigma.inv <- inverse(Sigma)

    # Set up mean and covariance matrix

    for(i in 1:T) {
        Mu[i] <- alpha
        for (p in 1:14){
            C[i,i,k]=sigma[p]*X[i,k]
        }
        Sigma[i,i] <- sum(C[i,i,])
        for(j in (i+1):T) {
            for (p in 1:14){

```

Kernel	DIC	Running time(minutes)
Squared exponential	137.6	150
OU process	135	120
Brownian motion	53	25

Table 1: Comparison of three kernel

```

        C[i,j,p]=sigma[p]*min(X[i,p],X[j,p])
      }
      Sigma[i,j] <- sum(C[i,j,])
      Sigma[j,i] <- Sigma[i,j]
    }
  }
  #Priors

  alpha ~ dnorm(0, 0.01)
  for (p in 1:14){
    sigma[p] ~ dunif(0.1, 2)
  }

}"
dat<- list(Y=Yr,T=100,X=Xr)
jags.param=c('alpha','sigma')
fitbm1 <- jags(data=dat, n.chains=4, inits=NULL,parameters=jags.param, n.iter=3000, n.burnin=1000,
              model.file=textConnection(gpbm1))

```

The Brownian motion Gaussian process takes about 25 minutes to run with DIC 53, which is much more computationally convenient. Hence I choose Brownian motion as kernel function with hyperparameter  $\sigma_p \sim U(1, 2)$

The comparison of 3 kernel function is shown as below:

Since the value of  $\sigma_p$  are very close to each other, I would simplify the model that all the predictor share same priors  $\sigma \sim U(1, 2)$  in 5-fold cross validation for computationally convenient.

```

library(R2jags)

## Loading required package: rjags
## Loading required package: coda
## Linked to JAGS 4.3.0
## Loaded modules: basemod,bugs
##
## Attaching package: 'R2jags'
## The following object is masked from 'package:coda':
##
##      traceplot

stroke=read.table("/Users/jing/Desktop/2020 winter/Stat 225/hw2/Stroke.csv",sep=',',header=T)
stroke$Gender<-as.numeric(stroke$Gender)-1
attach(stroke)

strokecv<-stroke[sample(nrow(stroke)),]
#Create 5 equally size folds

```

```
folds <- cut(seq(1,nrow(strokecv)),breaks=5,labels=FALSE)
```

```
auc2<- matrix(nrow=5,ncol=1)
library(pROC)
```

```
## Type 'citation("pROC")' for a citation.
```

```
##
```

```
## Attaching package: 'pROC'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
##      cov, smooth, var
```

```
tp1 <- Sys.time()
```

```
for (l in 1:5) {
```

```
  testIndexes <- which(folds==l,arr.ind=TRUE)
```

```
  testData <- stroke[testIndexes, ]
```

```
  trainData <- stroke[-testIndexes, ]
```

```
  Xr<-trainData[,2:15]
```

```
  Yr<-trainData[,16]
```

```
  Xt<-testData[,2:15]
```

```
  Yt<-testData[,16]
```

```
gpbm2 <- "model{
```

```
  # Likelihood
```

```
  for(i in 1:T){
```

```
    Y[i] ~ dbern(p[i])
```

```
    p[i]<-exp(u[i])/(1+exp(u[i]))
```

```
  }
```

```
  u~ dmnorm(Mu, Sigma.inv)
```

```
  Sigma.inv <- inverse(Sigma)
```

```
# Set up mean and covariance matrix
```

```
for(i in 1:T) {
```

```
  Mu[i] <- alpha
```

```
  Sigma[i,i] <-sum(X[i,])*pow(sigma, 2)
```

```
  for(j in (i+1):T) {
```

```
    for (p in 1:14){
```

```
      C[i,j,p]=min(X[i,p],X[j,p])
```

```
    }
```

```
    Sigma[i,j] <- pow(sigma, 2) * sum(C[i,j,])
```

```
    Sigma[j,i] <- Sigma[i,j]
```

```
  }
```

```
}
```

```
#Priors
```

```
alpha ~ dnorm(0, 0.01)
```

```
sigma ~ dunif(1, 2)
```

```

}"

dat<- list(Y=Yr,T=80,X=Xr)
jags.param=c('alpha','sigma')
fitbm2<- jags(data=dat, n.chains=4, inits=NULL,parameters=jags.param, n.iter=3000, n.burnin=1000,
              model.file=textConnection(gpbm2))
print(fitbm2)
alpha = fitbm2$BUGSoutput$sims.list$alpha
sigma=fitbm2$BUGSoutput$sims.list$sigma
#check the posterior distribution
par(mfrow = c(1,2))
hist(alpha,30)
hist(sigma,30)

# prediction
T_new = 20
Mu = rep(mean(alpha), 80)
Mu_new = rep(mean(alpha), 20)
Sigma=matrix(rep(0,80*80),nrow=80,ncol=80)
#Sigma_new = mean(sigma)^2 * exp( -mean(rho) * outer(t, t_new, '-')^2 )

#pred_mean = Mu_new + t(Sigma_new)%*%solve(Sigma, y - Mu)

C_function = function(x1,x2,k) {
  Sigma <- matrix(rep(0, length(x1)*length(x2)), nrow=length(x1))
  for (i in 1:nrow(Sigma)) {
    for (j in 1:ncol(Sigma)) {
      Sigma[i,j] <- mean((sigma)^2) *min(x1[i],x2[j])
    }
  }
  return(Sigma)
}

for (p in 1:14){
  Sigma=Sigma+C_function(Xr[,p],Xr[,p],p)
}

Sigma_new=matrix(rep(0,20*80),nrow=20,ncol=80)
for (p in 1:14){
  Sigma_new=Sigma_new+C_function(Xt[,p],Xr[,p],p)
}

for (i in 1:80){
  if (Yr[i]==1){Yr[i]=Yr[i]-0.0001}
  if (Yr[i]==0){Yr[i]=Yr[i]+0.0001}
}

# compute the posterior predictive distribution
pred_mean = Mu_new + Sigma_new%*%solve(Sigma, log(Yr/(1-Yr)) - Mu)

```

```

p=exp(pred_mean)/(1+exp(pred_mean))


auc2[l]<-auc(Yt, p)
#par(mfrow = c(1,1))
#print(roc(Yt, p,plot = T,levels=c("0", "1"), direction="<"))
}

## module glm loaded


## Compiling model graph
##   Resolving undeclared variables
##   Allocating nodes
## Graph information:
##   Observed stochastic nodes: 80
##   Unobserved stochastic nodes: 3
##   Total graph size: 43390
##
## Initializing model
##
## Inference for Bugs model at "4", fit using jags,
## 4 chains, each with 3000 iterations (first 1000 discarded), n.thin = 2
## n.sims = 4000 iterations saved
##      mu.vect sd.vect   2.5%   25%   50%   75%  97.5%  Rhat n.eff
## alpha      0.396   5.314 -9.854 -3.173  0.388  3.832 11.239 1.021   120
## sigma      1.048   0.049  1.001  1.013  1.031  1.065  1.185 1.040    93
## deviance   54.829   7.846 39.896 49.307 54.793 59.647 71.146 1.625     8
##
## For each parameter, n.eff is a crude measure of effective sample size,
## and Rhat is the potential scale reduction factor (at convergence, Rhat=1).
##
## DIC info (using the rule, pD = var(deviance)/2)
## pD = 18.0 and DIC = 72.8
## DIC is an estimate of expected predictive error (lower deviance is better).
## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

```

### Histogram of alpha




### Histogram of sigma




```
## Compiling model graph
##   Resolving undeclared variables
##   Allocating nodes
## Graph information:
##   Observed stochastic nodes: 80
##   Unobserved stochastic nodes: 3
##   Total graph size: 43438
##
## Initializing model
##
## Inference for Bugs model at "5", fit using jags,
## 4 chains, each with 3000 iterations (first 1000 discarded), n.thin = 2
## n.sims = 4000 iterations saved
##      mu.vect sd.vect   2.5%   25%   50%   75%  97.5%  Rhat n.eff
## alpha    -1.217  5.518 -12.112 -4.915 -1.147  2.526  9.377 1.025  110
## sigma     1.068  0.075  1.001  1.017  1.042  1.093  1.267 1.242   19
## deviance  56.811  7.845  41.197 51.567 57.202 62.335 71.055 1.976    6
##
## For each parameter, n.eff is a crude measure of effective sample size,
## and Rhat is the potential scale reduction factor (at convergence, Rhat=1).
##
## DIC info (using the rule, pD = var(deviance)/2)
## pD = 14.0 and DIC = 70.8
## DIC is an estimate of expected predictive error (lower deviance is better).
##
## Setting levels: control = 0, case = 1
## Setting direction: controls < cases
```



### Histogram of alpha




### Histogram of sigma




```
## Compiling model graph
##   Resolving undeclared variables
##   Allocating nodes
## Graph information:
##   Observed stochastic nodes: 80
##   Unobserved stochastic nodes: 3
##   Total graph size: 43353
##
## Initializing model
##
## Inference for Bugs model at "4", fit using jags,
## 4 chains, each with 3000 iterations (first 1000 discarded), n.thin = 2
## n.sims = 4000 iterations saved
##      mu.vect sd.vect   2.5%   25%   50%   75%  97.5%  Rhat n.eff
## alpha    -1.369   5.271 -11.734 -4.931 -1.360  2.170  8.758 1.011   240
## sigma     1.046   0.049   1.001  1.012  1.031  1.061  1.180 1.062    73
## deviance  58.377   9.089  41.529 51.902 58.396 64.679 76.043 1.635     8
##
## For each parameter, n.eff is a crude measure of effective sample size,
## and Rhat is the potential scale reduction factor (at convergence, Rhat=1).
##
## DIC info (using the rule, pD = var(deviance)/2)
## pD = 23.4 and DIC = 81.8
## DIC is an estimate of expected predictive error (lower deviance is better).
##
## Setting levels: control = 0, case = 1
## Setting direction: controls < cases
```

### Histogram of alpha




### Histogram of sigma




```
## Compiling model graph
##   Resolving undeclared variables
##   Allocating nodes
## Graph information:
##   Observed stochastic nodes: 80
##   Unobserved stochastic nodes: 3
##   Total graph size: 43469
##
## Initializing model
##
## Inference for Bugs model at "6", fit using jags,
## 4 chains, each with 3000 iterations (first 1000 discarded), n.thin = 2
## n.sims = 4000 iterations saved
##      mu.vect sd.vect   2.5%   25%   50%   75%  97.5%  Rhat n.eff
## alpha  -0.630  5.170 -10.780 -4.071 -0.766  2.809  9.481 1.009  290
## sigma   1.038  0.037  1.001  1.011  1.026  1.053  1.135 1.026  160
## deviance 54.683  8.963  37.239 48.407 55.232 61.265 71.030 1.486   10
##
## For each parameter, n.eff is a crude measure of effective sample size,
## and Rhat is the potential scale reduction factor (at convergence, Rhat=1).
##
## DIC info (using the rule, pD = var(deviance)/2)
## pD = 26.4 and DIC = 81.1
## DIC is an estimate of expected predictive error (lower deviance is better).
##
## Setting levels: control = 0, case = 1
## Setting direction: controls < cases
```

### Histogram of alpha




### Histogram of sigma




```
## Compiling model graph
##   Resolving undeclared variables
##   Allocating nodes
## Graph information:
##   Observed stochastic nodes: 80
##   Unobserved stochastic nodes: 3
##   Total graph size: 43373
##
## Initializing model
##
## Inference for Bugs model at "4", fit using jags,
## 4 chains, each with 3000 iterations (first 1000 discarded), n.thin = 2
## n.sims = 4000 iterations saved
##      mu.vect sd.vect   2.5%   25%   50%   75%  97.5%  Rhat n.eff
## alpha    -1.151   5.380 -11.715 -4.816 -1.218  2.314  9.553 1.030   87
## sigma     1.048   0.054  1.001  1.012  1.031  1.065  1.216 1.131   39
## deviance  54.994   9.045  41.547 47.996 53.157 60.949 74.697 1.960    6
##
## For each parameter, n.eff is a crude measure of effective sample size,
## and Rhat is the potential scale reduction factor (at convergence, Rhat=1).
##
## DIC info (using the rule, pD = var(deviance)/2)
## pD = 17.8 and DIC = 72.8
## DIC is an estimate of expected predictive error (lower deviance is better).
##
## Setting levels: control = 0, case = 1
## Setting direction: controls < cases
```

### Histogram of alpha



### Histogram of sigma



```
tp2 <- Sys.time()
tp2-tp1
```

```
## Time difference of 7.857199 mins
```

```
print(mean(auc2))
```

```
## [1] 0.8177529
```

The mean AUC of 5 fold cross validation in Gaussian process is 0.82 which share the similar results with logistic model I built in homework 2. Therefore, the final kernel I choose is

$$C_{ij} = \sigma^2 \sum_{p=1}^p \min(X_{ip}, X_{jp})$$

, where  $\sigma \sim U(1, 2)$ .